

Discussion Section 1

GOVT 6039

Georgy Tarasenko

What quantity would answer this question?

Does making local news free
make people better informed?

?

Write down one quantity you would report if you could observe the entire population — in words, or as a formula if you can manage one.

Three targets behind one sentence

A = 1 if a person was offered free access, = 0 otherwise

Y score on a local-politics quiz, from 0 to 20

Three targets behind one sentence

A = 1 if a person was offered free access, = 0 otherwise

Y score on a local-politics quiz, from 0 to 20

$$\mu = \underbrace{E[Y]}_{\text{average over everyone}}$$

average knowledge

Three targets behind one sentence

A = 1 if a person was offered free access, = 0 otherwise

Y score on a local-politics quiz, from 0 to 20

$\mu = \underbrace{E[Y]}_{\text{average over everyone}}$ average knowledge

$p_{15} = \underbrace{P(Y \geq 15)}_{\text{fraction above a cutoff}}$ share classified as informed

Three targets behind one sentence

A = 1 if a person was offered free access, = 0 otherwise

Y = score on a local-politics quiz, from 0 to 20

$\mu = \underbrace{E[Y]}_{\text{average over everyone}}$ average knowledge

$p_{15} = \underbrace{P(Y \geq 15)}_{\text{fraction above a cutoff}}$ share classified as informed

$\Delta = \underbrace{E[Y | A=1]}_{\text{mean among those offered}} - \underbrace{E[Y | A=0]}_{\text{mean among those not offered}}$ the free-access gap

Three targets behind one sentence

A = 1 if a person was offered free access, = 0 otherwise

Y = score on a local-politics quiz, from 0 to 20

$\mu = \underbrace{E[Y]}_{\text{average over everyone}}$ average knowledge


$p_{15} = \underbrace{P(Y \geq 15)}_{\text{fraction above a cutoff}}$ share classified as informed

$\Delta = \underbrace{E[Y | A=1]}_{\text{mean among those offered}} - \underbrace{E[Y | A=0]}_{\text{mean among those not offered}}$ the free-access gap

Nothing in the data tells you which of the three to report. That choice comes from the research question.


Same intervention, different summaries

Suppose free access raises the scores of people who were already close to the cutoff, and leaves everyone else where they were.



Same intervention, different summaries

Suppose free access raises the scores of people who were already close to the cutoff, and leaves everyone else where they were.



▲ marks the mean


red line is the cutoff 15

shaded is the region $Y \geq 15$

μ : 11.0 $\rightarrow$ 12.4

Same intervention, different summaries

Suppose free access raises the scores of people who were already close to the cutoff, and leaves everyone else where they were.



▲ marks the mean

red line is the cutoff 15


shaded is the region $Y \geq 15$

μ : 11.0 $\rightarrow$ 12.4

p_{15} : 0.10 $\rightarrow$ 0.44

Same intervention, different summaries

Suppose free access raises the scores of people who were already close to the cutoff, and leaves everyone else where they were.



▲ marks the mean

red line is the cutoff 15

shaded is the region $Y \geq 15$

μ : 11.0 $\rightarrow$ 12.4

p_{15} : 0.10 $\rightarrow$ 0.44

The intervention is identical in both panels. How big the change looks depends on which summary you report.

Target, rule, or result?

H is hours per week a person spends with local news. D is the dataset we happen to have.

- 1 **estimand** — a feature of the population, fixed before any data
- 2 **estimator** — a procedure, a rule you could run on any dataset
- 3 **estimate** — the number that procedure returns on *this* dataset

Target, rule, or result?

H is hours per week a person spends with local news. D is the dataset we happen to have.

- 1 **estimand** — a feature of the population, fixed before any data
- 2 **estimator** — a procedure, a rule you could run on any dataset
- 3 **estimate** — the number that procedure returns on *this* dataset

$$f(6) = E[Y \mid H = 6]$$

Target, rule, or result?

H is hours per week a person spends with local news. D is the dataset we happen to have.

- 1 **estimand** — a feature of the population, fixed before any data
- 2 **estimator** — a procedure, a rule you could run on any dataset
- 3 **estimate** — the number that procedure returns on *this* dataset

$$f(6) = E[Y \mid H = 6]$$

fit a cubic by minimising the sum of squared residuals

Target, rule, or result?

H is hours per week a person spends with local news. D is the dataset we happen to have.

- 1 **estimand** — a feature of the population, fixed before any data
- 2 **estimator** — a procedure, a rule you could run on any dataset
- 3 **estimate** — the number that procedure returns on *this* dataset

$$f(6) = E[Y \mid H = 6]$$

fit a cubic by minimising the sum of squared residuals

$$\hat{f}_D(6) = 14.1$$

Target, rule, or result?

H is hours per week a person spends with local news. D is the dataset we happen to have.

- 1 **estimand** — a feature of the population, fixed before any data
- 2 **estimator** — a procedure, a rule you could run on any dataset
- 3 **estimate** — the number that procedure returns on *this* dataset

$$f(6) = E[Y \mid H = 6]$$

fit a cubic by minimising the sum of squared residuals

$$\hat{f}_D(6) = 14.1$$

? Is each of the three objects above a **1**, a **2**, or a **3**? Hold up fingers.

Target, rule, or result?

H is hours per week a person spends with local news. D is the dataset we happen to have.

- 1 **estimand** — a feature of the population, fixed before any data
- 2 **estimator** — a procedure, a rule you could run on any dataset
- 3 **estimate** — the number that procedure returns on *this* dataset

$$f(6) = E[Y \mid H = 6]$$

fit a cubic by minimising the sum of squared residuals

$$\hat{f}_D(6) = 14.1$$

? Is each of the three objects above a **1**, a **2**, or a **3**? Hold up fingers.

$f(6)$ is **1** · the cubic procedure is **2** · 14.1 is **3**

The pipeline



The pipeline



?

We draw a second sample from the same population and run the same procedure. Which of the four objects take different values?

The pipeline



$P \longrightarrow D$ the population also generates the dataset

?

We draw a second sample from the same population and run the same procedure. Which of the four objects take different values?

The pipeline



$P \longrightarrow D$ the population also generates the dataset

?

We draw a second sample from the same population and run the same procedure. Which of the four objects take different values?

D and $\hat{\theta}$ take new values. θ does not move. $\mathcal{A}$ moves only if the analyst rewrites it.

Which edge broke, and what would repair it?

Three studies of the same question: does free access to local news raise political knowledge?

- A** Two independent samples of 300 people give free-access gaps of 2.1 and 0.4 points.
- B** The target is the gap among people who already follow local politics. The fitted model has no interaction term, so it reports one gap for followers and non-followers alike.
- C** The policy question is whether people cross a threshold of being informed. All the study reports is the mean quiz score.

? For each case, name the edge that broke and one repair.

Different breaks need different repairs

Failure

What went wrong

Appropriate response

Different breaks need different repairs

Failure	What went wrong	Appropriate response
question $\rightarrow$ estimand	wrong target	rewrite θ

Different breaks need different repairs

Failure	What went wrong	Appropriate response
question $\rightarrow$ estimand	wrong target	rewrite θ
estimand $\rightarrow$ estimator	the rule cannot recover or represent the target	change the model, the procedure, or the assumptions

Different breaks need different repairs

Failure	What went wrong	Appropriate response
question → estimand	wrong target	rewrite θ
estimand → estimator	the rule cannot recover or represent the target	change the model, the procedure, or the assumptions
estimator → estimate	sampling variability, or problematic data	more or better data; quantify uncertainty

Different breaks need different repairs

Failure	What went wrong	Appropriate response
question → estimand	wrong target	rewrite θ
estimand → estimator	the rule cannot recover or represent the target	change the model, the procedure, or the assumptions
estimator → estimate	sampling variability, or problematic data	more or better data; quantify uncertainty
estimate → claim	the sentence exceeds the evidence	narrow the claim, or validate further

Different breaks need different repairs

Failure	What went wrong	Appropriate response
question → estimand	wrong target	rewrite θ
estimand → estimator	the rule cannot recover or represent the target	change the model, the procedure, or the assumptions
estimator → estimate	sampling variability, or problematic data	more or better data; quantify uncertainty
estimate → claim	the sentence exceeds the evidence	narrow the claim, or validate further

C is the first row, B the second, A the third.

Different breaks need different repairs

Failure	What went wrong	Appropriate response
question → estimand	wrong target	rewrite θ
estimand → estimator	the rule cannot recover or represent the target	change the model, the procedure, or the assumptions
estimator → estimate	sampling variability, or problematic data	more or better data; quantify uncertainty
estimate → claim	the sentence exceeds the evidence	narrow the claim, or validate further

C is the first row, B the second, A the third.

Only the third row is fixed by collecting more data. The other three need somebody to change the target, the model, or the sentence.

Four cells, three questions

A = 1 if offered free access. Randomly assigned.

Z = 1 if the person already follows local politics closely. Observed, not assigned.

	$A = 0$	$A = 1$
$Z = 1$	18	19
$Z = 0$	10	15

each cell is $E[Y \mid A, Z]$

Y = quiz score. The four numbers below are population means, known exactly.

π = $P(Z=1) = 0.5$, the share of close followers.

- 1 **Level.** Expected knowledge among people offered access *and* already following.
- 2 **Heterogeneity.** How much smaller is the access gap among followers than among non-followers?
- 3 **Population average.** The access gap for one person drawn at random.

? For each: a number, an expression in the four cells, and one sentence saying what it means.

Four cells, three answers

	$A = 0$	$A = 1$
$Z = 1$	18	19
$Z = 0$	10	15

Four cells, three answers

	$A = 0$	$A = 1$
$Z = 1$	18	19
$Z = 0$	10	15

$$\theta_L = E[Y \mid A=1, Z=1] = 19$$

one cell, read off directly

Four cells, three answers

	$A = 0$		$A = 1$
$Z = 1$	18	$\Delta_1 = 1$	19
$Z = 0$	10	$\Delta_0 = 5$	15

$\theta_L = E[Y \mid A=1, Z=1] = 19$ one cell, read off directly

$\Delta_1 = 19 - 18 = 1$ and $\Delta_0 = 15 - 10 = 5$ the gap inside each row

Four cells, three answers

	$A = 0$	$A = 1$
$Z = 1$	18	19
$Z = 0$	10	15

Diagram illustrating the four cells of a 2x2 factorial design. The columns represent treatment A ($A=0$ and $A=1$), and the rows represent treatment Z ($Z=1$ and $Z=0$). The values in the cells are 18, 19, 10, and 15. The differences between the cells in each row are labeled $\Delta_1 = 1$ (for $Z=1$) and $\Delta_0 = 5$ (for $Z=0$).

$$\theta_L = E[Y \mid A=1, Z=1] = 19 \quad \text{one cell, read off directly}$$

$$\Delta_1 = 19 - 18 = 1 \quad \text{and} \quad \Delta_0 = 15 - 10 = 5 \quad \text{the gap inside each row}$$

$$\theta_H = \Delta_1 - \Delta_0 = 1 - 5 = -4 \quad \text{how much smaller the gap is for followers}$$

Four cells, three answers

	$A = 0$		$A = 1$	
$Z = 1$	18	$\Delta_1 = 1$	19	$\times \pi$
$Z = 0$	10	$\Delta_0 = 5$	15	$\times (1 - \pi)$

$$\theta_L = E[Y \mid A=1, Z=1] = 19 \quad \text{one cell, read off directly}$$

$$\Delta_1 = 19 - 18 = 1 \quad \text{and} \quad \Delta_0 = 15 - 10 = 5 \quad \text{the gap inside each row}$$

$$\theta_H = \Delta_1 - \Delta_0 = 1 - 5 = -4 \quad \text{how much smaller the gap is for followers}$$

$$\theta_A = (1 - \pi)\Delta_0 + \pi\Delta_1 = 0.5(5) + 0.5(1) = 3 \quad \text{each gap weighted by how common its row is}$$

The same four cells, written as coefficients

$$f(a, z) = E[Y \mid A = a, Z = z] = \underbrace{\beta_0}_{\text{cell } (0,0)} + \underbrace{\beta_1 a}_{\text{shift from } A} + \underbrace{\beta_2 z}_{\text{shift from } Z} + \underbrace{\beta_3 az}_{\text{extra shift when both}}$$

$$\beta_0 = 10 \quad \beta_1 = 5 \quad \beta_2 = 8 \quad \beta_3 = -4$$

The same four cells, written as coefficients

$$f(a, z) = E[Y \mid A = a, Z = z] = \underbrace{\beta_0}_{\text{cell } (0,0)} + \underbrace{\beta_1 a}_{\text{shift from } A} + \underbrace{\beta_2 z}_{\text{shift from } Z} + \underbrace{\beta_3 az}_{\text{extra shift when both}}$$

$$\beta_0 = 10 \quad \beta_1 = 5 \quad \beta_2 = 8 \quad \beta_3 = -4$$

$$\theta_L = \beta_0 + \beta_1 + \beta_2 + \beta_3$$

$$\theta_H = \beta_3$$

$$\theta_A = \beta_1 + \pi\beta_3$$

The same four cells, written as coefficients

$$f(a, z) = E[Y \mid A = a, Z = z] = \underbrace{\beta_0}_{\text{cell } (0,0)} + \underbrace{\beta_1 a}_{\text{shift from } A} + \underbrace{\beta_2 z}_{\text{shift from } Z} + \underbrace{\beta_3 az}_{\text{extra shift when both}}$$

$$\beta_0 = 10 \quad \beta_1 = 5 \quad \beta_2 = 8 \quad \beta_3 = -4$$

$$\theta_L = \beta_0 + \beta_1 + \beta_2 + \beta_3$$

one cell

$$\theta_H = \beta_3$$

difference between gaps

$$\theta_A = \beta_1 + \pi\beta_3$$

weighted average of gaps

The same four cells, written as coefficients

$$f(a, z) = E[Y \mid A = a, Z = z] = \underbrace{\beta_0}_{\text{cell } (0,0)} + \underbrace{\beta_1 a}_{\text{shift from } A} + \underbrace{\beta_2 z}_{\text{shift from } Z} + \underbrace{\beta_3 az}_{\text{extra shift when both}}$$

$$\beta_0 = 10 \quad \beta_1 = 5 \quad \beta_2 = 8 \quad \beta_3 = -4$$

$$\theta_L = \beta_0 + \beta_1 + \beta_2 + \beta_3$$

one cell

$$\theta_H = \beta_3$$

difference between gaps

$$\theta_A = \beta_1 + \pi\beta_3$$

weighted average of gaps

The average in θ_A runs over the two *gaps* Δ_0 and Δ_1 , with θ_H playing no part in it.

The short model reports 3. Which target is that?

The model has no interaction term, so one A -gap is forced onto both rows.

$$\hat{f}_{\text{short}}(a, z) = \hat{\alpha}_0 + \hat{\alpha}_1 a + \hat{\alpha}_2 z$$

$$\hat{\alpha}_1 = 3$$



one gap, forced onto both rows

? Which of the three targets does the number 3 answer?

- A** the gap among non-followers, $\Delta_0 = 5$
- B** the gap among followers, $\Delta_1 = 1$
- C** the population-average gap, $\theta_A = 3$
- D** all three at once

The answer is C

The answer is C

$$\theta_A = (1 - \pi)\Delta_0 + \pi\Delta_1 = 0.5(5) + 0.5(1) = 3$$

$$\theta_A = (1 - \pi)\Delta_0 + \pi\Delta_1 = 0.5(5) + 0.5(1) = 3$$

The subgroup gaps are $\Delta_0 = 5$ and $\Delta_1 = 1$, so 3 is not either of them.

$$\theta_A = (1 - \pi)\Delta_0 + \pi\Delta_1 = 0.5(5) + 0.5(1) = 3$$

The subgroup gaps are $\Delta_0 = 5$ and $\Delta_1 = 1$, so 3 is not either of them.

The model returns the right number for one target even though it has no way to represent either subgroup gap.


$$\theta_A = (1 - \pi)\Delta_0 + \pi\Delta_1 = 0.5(5) + 0.5(1) = 3$$

The subgroup gaps are $\Delta_0 = 5$ and $\Delta_1 = 1$, so 3 is not either of them.

The model returns the right number for one target even though it has no way to represent either subgroup gap.


Adding Z to the model does not let the A -gap vary with Z .

Infinite data cannot escape the model class



dotted black: the truth f · solid red: best member of the class · thin red: one finite sample

Infinite data cannot escape the model class



dotted black: the truth f · solid red: best member of the class · thin red: one finite sample

?

Which of the two gaps survives when the sample becomes infinite?

Two different gaps

$$\underbrace{f_{\mathcal{F}}^*}_{\text{best member of the class}} = \arg \min_{g \in \mathcal{F}} \underbrace{E[(f(H) - g(H))^2]}_{\text{average squared distance to the truth, over the population}}$$

Two different gaps

$$\underbrace{f_{\mathcal{F}}^*}_{\text{best member of the class}} = \arg \min_{g \in \mathcal{F}} \underbrace{E[(f(H) - g(H))^2]}_{\text{average squared distance to the truth, over the population}}$$

stays $f - f_{\mathcal{F}}^*$ the class cannot bend any further. This gap has nothing to do with n .

Two different gaps

$$\underbrace{f_{\mathcal{F}}^*}_{\text{best member of the class}} = \arg \min_{g \in \mathcal{F}} \underbrace{E[(f(H) - g(H))^2]}_{\text{average squared distance to the truth, over the population}}$$

stays $f - f_{\mathcal{F}}^*$ the class cannot bend any further. This gap has nothing to do with n .

shrinks $f_{\mathcal{F}}^* - \hat{f}_D$ the sample is finite. This gap falls as n grows.

Two different gaps

$$\underbrace{f_{\mathcal{F}}^*}_{\text{best member of the class}} = \arg \min_{g \in \mathcal{F}} \underbrace{E[(f(H) - g(H))^2]}_{\text{average squared distance to the truth, over the population}}$$

stays $f - f_{\mathcal{F}}^*$ the class cannot bend any further. This gap has nothing to do with n .


shrinks $f_{\mathcal{F}}^* - \hat{f}_D$ the sample is finite. This gap falls as n grows.

Least squares only ever sees your sample, so what it returns is the best fit to that sample. No dataset will hand you the solid red curve.

Offering access can change knowledge only through reading

In Trexler's experiment the researcher controls only the offer. Whether anyone reads anything, and what they learn from it, follows on from that.

randomly assigned




$$\Delta_M = \underbrace{E[M \mid A=1] - E[M \mid A=0]}_{\text{did the offer change reading?}}$$

$$\Delta_Y = \underbrace{E[Y \mid A=1] - E[Y \mid A=0]}_{\text{did the offer change knowledge?}}$$


Both group people by the offer they received rather than by what they did with it. These are intent-to-treat quantities.

Did the offer change reading?



Adapted from Trexler (2026), *JOP* 88(3), figure 2. Each row is one outcome. Dot is the ITT estimate, bar is the 95% interval, dashed line is zero.


Did the offer change reading?



Adapted from Trexler (2026), *JOP* 88(3), figure 2. Each row is one outcome. Dot is the ITT estimate, bar is the 95% interval, dashed line is zero.


Newsletter recall goes up by about 0.33, while both measures of actual consumption sit on zero.

Did it change knowledge?



Adapted from Trexler (2026), *JOP* 88(3), figure 3, knowledge block. Horizontal axis runs from -0.2 to 0.2 .

Did it change knowledge?



Adapted from Trexler (2026), *JOP* 88(3), figure 3, knowledge block. Horizontal axis runs from -0.2 to 0.2 .

All three knowledge estimates overlap zero.

“Reading local news does not make people better informed.”

?

Rewrite this so the treatment, the outcome, the population and the design all match what was measured.

“Reading local news does not make people better informed.”

?

Rewrite this so the treatment, the outcome, the population and the design all match what was measured.

Offering two months of free digital access raised newsletter recall but did not measurably raise reported news consumption or political knowledge among the registered non-subscribers who joined the study.

Rewrite the claim

“Reading local news does not make people better informed.”

?


Rewrite this so the treatment, the outcome, the population and the design all match what was measured.

Offering two months of free digital access raised newsletter recall but did not measurably raise reported news consumption or political knowledge among the registered non-subscribers who joined the study.

The offer is what got randomized, so the offer is what this experiment measures.

The same weather history, two different treatments

Quoß and Rudolph ask whether extreme weather shifts climate policy preferences. Someone must first decide what counts as extreme.




Adapted from Quoß & Rudolph (2026), *JOP* 88(3), figure 4. Each dot is one respondent. A: IPCC definition. B: MeteoSwiss definition.

$$\underbrace{W_i^{(\text{IPCC})} \neq W_i^{(\text{MeteoSwiss})}}_{\text{the same person gets a different exposure value}} \quad \theta^{(k)} = T(P; W^{(k)})$$

? Two rival estimators of one estimand, or two different estimands?

One definition, or thirty-four?


The authors run the same design, on the same people, with the same outcome, under 34 defensible definitions of extreme weather.



Adapted from Quoß & Rudolph (2026), *JOP* 88(3), figure 6. Each dot is the coefficient from one definition, sorted low to high. Bars are 95% and 90% intervals. Filled dots are significant at 5%.

One definition, or thirty-four?

The authors run the same design, on the same people, with the same outcome, under 34 defensible definitions of extreme weather.




Adapted from Quoß & Rudolph (2026), *JOP* 88(3), figure 6. Each dot is the coefficient from one definition, sorted low to high. Bars are 95% and 90% intervals. Filled dots are significant at 5%.

Only $W^{(k)}$ changes across the 34 rows.

One definition, or thirty-four?

The authors run the same design, on the same people, with the same outcome, under 34 defensible definitions of extreme weather.



Adapted from Quoß & Rudolph (2026), *JOP* 88(3), figure 6. Each dot is the coefficient from one definition, sorted low to high. Bars are 95% and 90% intervals. Filled dots are significant at 5%.

Only $W^{(k)}$ changes across the 34 rows.

Any one of these definitions could be written up as a signed effect, but across all 34 most of the estimates sit on zero.

Reducible and irreducible, at one point

Take one fitted curve $\hat{f}_D$ and one value x . Draw a fresh outcome Y at that x and score the prediction.


$$\begin{aligned} R_D(x) &= \underbrace{E[(Y - \hat{f}_D(x))^2 \mid X = x, D]}_{\text{average squared miss on a fresh draw, with } D \text{ held fixed}} \\ &= \underbrace{[f(x) - \hat{f}_D(x)]^2}_{\text{reducible: how far this curve sits from the truth}} + \underbrace{\sigma^2}_{\text{irreducible: noise around } f} \end{aligned}$$

At one point x_0 , three fitted curves score $R_D^A(x_0) = 89$ $R_D^B(x_0) = 34$ $R_D^C(x_0) = 25$, and we are told $\hat{f}_C(x_0) = f(x_0)$.


- 1 Find σ^2 .
- 2 Find $|f(x_0) - \hat{f}_D(x_0)|$ for A and B .
- 3 Does curve B sit above or below $f(x_0)$?
- 4 Could a fourth curve score $R_D(x_0) = 20$?

?

Reducible and irreducible, solved




Reducible and irreducible, solved



Curve C has no reducible part, so its whole score is the floor: $\sigma^2 = 25$. Then $\sqrt{89 - 25} = 8$ and $\sqrt{34 - 25} = 3$.


Reducible and irreducible, solved



Curve C has no reducible part, so its whole score is the floor: $\sigma^2 = 25$. Then $\sqrt{89 - 25} = 8$ and $\sqrt{34 - 25} = 3$.

Above or below is **not identified**: the formula holds $[f - \hat{f}]^2$, and squaring removes the sign.

Reducible and irreducible, solved



Curve C has no reducible part, so its whole score is the floor: $\sigma^2 = 25$. Then $\sqrt{89 - 25} = 8$ and $\sqrt{34 - 25} = 3$.

Above or below is **not identified**: the formula holds $[f - \hat{f}]^2$, and squaring removes the sign.

Since $[f(x_0) - \hat{f}_D(x_0)]^2 \geq 0$, every curve scores $R_D(x_0) \geq 25$. A score of 20 cannot be this quantity.

Complete the pipeline

My research question is _____.

$\theta =$ _____

I would estimate it using _____.

The strongest claim this design could support is _____.

If the second line contains $\hat{\beta}$, the name of a regression, or “run OLS”, you have written an estimator rather than an estimand.

Trexler, A. (2026). "The Minimal Effects of Making Local News Free: Evidence from a Field Experiment." *Journal of Politics* 88(3). Figures 2 and 3.

Quoß, F. & Rudolph, L. (2026). "The Hidden Universe of Operationalizing Natural Experiments: Navigating Extreme Weather." *Journal of Politics* 88(3). Figures 4 and 6.

Lundberg, I., Johnson, R. & Stewart, B. M. (2021). "What Is Your Estimand?" *American Sociological Review* 86(3).

James, G., Witten, D., Hastie, T. & Tibshirani, R. (2021). *An Introduction to Statistical Learning*, 2nd ed., section 2.1.